

# Homework 1 Solutions

January 27, 2024

See also the accompanying Julia notebook for computational solutions.

## Problem 1 (5 points)

Start reading the draft course notes (linked from <https://github.com/mitmath/matrixcalc/>). Find a place that you found confusing, and write a paragraph explaining the source of your confusion and (ideally) suggesting a possible improvement.

(Any other corrections/comments are welcome, too.)

### Solution:

Student-dependent, but full marks if clearly written and explained.

## Problem 2 (5 points)

A directional derivative of  $f(x)$  in a direction  $v$  is sometimes described as the derivative  $\left. \frac{d}{d\alpha} f(x + \alpha v) \right|_{\alpha=0}$ , where  $\alpha \in \mathbb{R}$  is a scalar; that is, it is  $g'(0)$  for  $g(\alpha) = f(x + \alpha v)$ . If  $f(x)$  is a function from some input vector space  $x \in X$  to some output vector space  $f(x) \in Y$  with a derivative  $f'(x)$  as defined in class, apply the chain rule to obtain this  $g'(0)$  (for some  $v \in X$ ) in terms of  $f'$ .

### Solution:

Let  $h(\alpha) = x + \alpha v$ . We have  $g(\alpha) = f(h(\alpha))$ , so by the chain rule we must have  $g'(0)d\alpha = f'(h(0))[h'(0)d\alpha] = f'(x)[h'(0)]d\alpha$  (the scalar  $d\alpha$  can be pulled out by linearity). But  $dh = h'(\alpha)[d\alpha] = d\alpha v$ , so  $h'$  is simply the linear operator that multiplies  $d\alpha$  by  $v$ , i.e.  $h' = v$ . So, we have:

$$g'(0) = f'(x)[h'(0)] = \boxed{f'(x)[v]},$$

which is the same as the “directional derivative” defined in class.

For example, if  $x$  is a column vector and  $f(x)$  is a scalar, so that  $f' = (\nabla f)^T$ , this reduces to  $(\nabla f)^T v$  (the dot product of  $v$  with the gradient), which is another formula you may have seen.

## Problem 3 (3+3+3+3+3 points)

Find the derivatives  $f'$  of the following functions. If  $f$  maps column vectors to scalars, give  $\nabla f$  (so that  $f'(x)[dx] = (\nabla f)^T dx$  as in our definition of the gradient), and if  $f$  maps column vectors to column vectors gives the Jacobian matrix. Otherwise, simply write down  $f'$  as a linear operation.

1.  $f(x) = \|x\| = \sqrt{x^T x}$  for  $x \in \mathbb{R}^m$ .
2.  $f(x) = \frac{x^T (A + \|x\|^2 I) x}{x^T x}$  for  $x \in \mathbb{R}^m$ ,  $A$  being a constant  $m \times m$  matrix, and  $I$  being the  $m \times m$  identity matrix.
3.  $f(A) = A^{-2}$  where  $A$  is an  $m \times m$  matrix.
4.  $f(A) = (\text{trace } A)^9$  where  $A$  is an  $m \times m$  matrix.
5.  $f(x) = A(x .* x)$  where  $A$  is an  $m \times n$  matrix,  $x \in \mathbb{R}^n$ , and  $.*$  denotes *elementwise* multiplication (also called a Hadamard product) in Julia/Matlab notation.

**Solution:**

1. By the chain rule,  $df = \sqrt{x^T x} = \frac{d(x^T x)}{2\sqrt{x^T x}}$  (using familiar single-variable derivative of square root). From class (via the product rule),  $d(x^T x) = 2x^T dx$ , so this yields:

$$df = f'(x)[dx] = \frac{2x^T dx}{2\sqrt{x^T x}} = \underbrace{\frac{x^T}{\|x\|}}_{f'=(\nabla f)^T} dx$$

giving  $\boxed{\nabla f = x/\|x\|}$ .

2. We can use the quotient rule  $f = g/h \implies df = (h dg - g dh)/h^2 = (dg - f dh)/h$  just as in ordinary single-variable calculus (derivation from the product rule still works:  $dg = d(fh) = df h + f dh$ ). Here, the denominator  $h = x^T x$  so  $dh = 2x^T dx$  (from class, by the product rule). And the numerator is  $g = x^T (A + \|x\|^2 I)x$ . Applying the product rule to this, and using the result from class for  $d(x^T Ax) = x^T (A + A^T)dx$  along with  $\|x\|^2 = x^T x = h \implies d(\|x\|^2) = 2x^T dx$  again, we obtain for the numerator:

$$dg = x^T (A + A^T + 2(x^T x)I) dx + \underbrace{x^T (2(x^T dx)I)x}_{=2x^T x(x^T dx)=x^T (2x^T x)I} dx = x^T (A + A^T + 4(x^T x)I) dx,$$

where we have used the fact that multiplication by scalars like  $x^T dx$  commute with all matrix/vector operations. So, we obtain an overall derivative  $f'$  of

$$df = f'(x)[dx] = \frac{x^T (A + A^T + 4(x^T x)I) dx - f(x)(2x^T)dx}{x^T x} = (\nabla f)^T dx,$$

giving

$$\boxed{\nabla f = \frac{(A + A^T + 4(x^T x)I)x - 2f(x)x}{x^T x}}. \quad (1)$$

(This could be written in a few other ways)

3. By the product rule,  $d(A^{-2}) = d(A^{-1})A^{-1} + A^{-1}d(A^{-1})$ , and since we know  $d(A^{-1}) = -A^{-1}dA A^{-1}$  from class, we immediately obtain:

$$df = f'(A)[dA] = d(A^{-2}) = \boxed{-A^{-1}dA A^{-2} - A^{-2}dA A^{-1}},$$

which is a linear operation on  $dA$ .

4. Trace is a linear operation so  $d \text{trace } A = \text{trace } dA$ . Since the trace is a scalar  $\alpha$ , the familiar single-variable power-rule applies  $d(\alpha^9) = 9\alpha^8 d\alpha$ . Putting these together (chain rule), we obtain:

$$df = f'(A)[dA] = \boxed{9(\text{trace } A)^8 \text{trace } dA},$$

which is a linear operation on  $dA$ .

(We could turn this into a matrix gradient  $\nabla f = 9(\text{trace } A)^8 I$  using the Frobenius inner product, but you weren't required to do this.)

5.  $f(x) = A(x .* x)$  where  $A$  is an  $m \times n$  matrix,  $x \in \mathbb{R}^n$ , and  $.*$  denotes *elementwise* multiplication (also called a Hadamard product) in Julia/Matlab notation.

It is easy to derive an elementwise product rule  $d(x .* x) = (dx .* x) + (x .* dx) = 2x .* dx$  by applying the ordinary product rule to each element and noting that  $.*$  is commutative. Then, from the product rule, we have:

$$df = f'(x)[dx] = A d(x .* x) = 2A(x .* dx).$$

But since this is a function that maps column vectors to column vectors, we should be able to write this as a Jacobian matrix multiplying  $dx$ . How? Simply think about what these operations do to  $dx$ : we first multiply each element of  $dx$  by the corresponding element of  $x$ , and then we multiply by  $A$  (the dot product of each row of  $A$  by this vector). But that is equivalent to *scaling each column of  $A$*  by the corresponding element of  $x$ , and then multiplying the resulting matrix by  $dx$  in the ordinary way. In Julia's "broadcast" notation, this operation of scaling each column of  $A$  is simply  $2A .* x^T$ , so we can write

$$df = 2(A .* x^T)dx$$

and the Jacobian is  $\boxed{J = 2A .* x^T}$ . Equivalently, the entries of the Jacobian are  $\boxed{J_{ij} = 2A_{ij}x_j}$ . (Any clear notation/description is acceptable here.)

## Problem 4 (5 points)

Suppose that  $f(t) = A(t)$  is a function that maps scalars  $t \in \mathbb{R}$  to  $m \times n$  matrices  $A(t)$ . For example,  $A(t) = \begin{pmatrix} \sin(t) & 0 & \cos(t) \\ t & t^2 & t^3 \end{pmatrix}$ . Explain why  $f'(t)[dt]$ , following our general definition, must simply correspond to taking the ordinary single-variable calculus derivative of each element of  $A(t)$  (the “elementwise” derivative) and multiplying it by the scalar  $dt$ . That is,  $f'(t) = A'(t)$  is the elementwise derivative.

### Solution:

From our general definition,  $df = f'(t)dt = dA = A(t+dt) - A(t)$ . But because matrix subtraction is done elementwise, this is a matrix whose entries are  $(dA)_{ij} = d(A_{ij}) = A_{ij}(t+dt) - A_{ij}(t) = A'_{ij}(t)dt$ .  $A_{ij}(t)$  is a single-variable function mapping  $\mathbb{R}$  to  $\mathbb{R}$ , however, so ordinary single-variable calculus applies to the derivative  $A'_{ij}(t)$ . Hence  $f'(t) = A'(t)$  is simply the elementwise derivative.

## Problem 5 (3+3+6+3 points)

If you are not familiar with 2d convolution operations (or even if you are), watch ( a little of the start of ) the YouTube video <https://www.youtube.com/watch?v=yb2tPt0QVPY> that explains them.

1. The (linear) convolution of an  $m \times n$  array (“matrix”) with the  $3 \times 3$  Sobel kernel  $\begin{pmatrix} 1 \\ 2 \\ 1 \end{pmatrix} \begin{pmatrix} -1 & 0 & 1 \end{pmatrix} = \begin{pmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{pmatrix}$  results in a \_\_\_\_\_ by \_\_\_\_\_ array?
2. Explain why this map from arrays to arrays is a linear operation.
3. Let  $X$  be a  $2024 \times 2024$  array, and  $Y$  be the result of convolving  $X$  with Sobel. Describe the matrix  $M$  that satisfies:

$$\text{vec}(Y) = M \text{vec}(X).$$

What is the size of  $M$ ? Express  $M$  in terms of Kronecker products of much smaller matrices (hint: 2d convolution with this Sobel kernel is “separable” into 1d convolutions acting on the rows and columns of  $X$ , and you can express these as matrices multiplying  $X$  on the \_\_\_\_\_ and \_\_\_\_\_, respectively).

4. Convolutions like this are very common linear operations, and yet they are not normally implemented by constructing an explicit matrix then multiplying it by a vector (even for 1d convolutions, much less 2d), no matter what you may have learned in linear algebra classes. Why is that?

### Solution:

1. As defined in the linked video, a convolution with a  $3 \times 3$  kernel by default shrinks the image by 1 in each direction from the edges—otherwise, you would need to define a “boundary condition” describing what to do for values beyond the edge of the image. So, the result should be a  $(m-2) \times (n-2)$  array.
2. The definition of convolution  $Y = f(X)$  is that each element of  $Y$  is a linear combination of the entries of  $X$ , which is obviously linear. In this particular case,  $Y_{ij} = -X_{i-1,j-1} + X_{i-1,j+1} - 2X_{i,j-1} + 2X_{i,j+1} - X_{i+1,j-1} + X_{i+1,j+1}$ , which is linear in  $X$ .
3. The input is a  $2024 \times 2024$  array (“matrix”)  $X$ , and hence from above the output is a  $2022 \times 2022$  array  $Y$ . Hence  $\text{vec}(X)$  has  $2024^2$  entries and  $\text{vec}(Y) = M \text{vec}(X)$  has  $2022^2$  entries, so  $M$  must be  $2022^2 \times 2024^2 = 4088484 \times 4096576$ , a huge matrix.

More explicitly, we can write this  $M$  in terms of Kronecker products of two  $2022 \times 2024$  matrices, because this Sobel kernel is *separable* into a convolution of each *column* of  $X$  with  $[1, 2, 1]$  (an “averaging” kernel, without the  $1/4$  normalization), and of each *row* of  $X$  with  $[-1, 0, 1]$  (a “difference” kernel). That is, each column is multiplied by the  $2022 \times 2024$  convolution matrix:

$$C = \begin{pmatrix} 1 & 2 & 1 & & \\ & 1 & 2 & 1 & \\ & & 1 & 2 & 1 \\ & & & \ddots & \ddots & \ddots \end{pmatrix}$$

To multiply *each column* of  $X$  by  $C$ , we simply do  $CX$ . Similarly, to convolve each row with  $[-1, 0, 1]$  corresponds to  $XX^T = (RX^T)^T$ , i.e. convolving each column of  $X^T$  (= rows of  $X$ ) with the  $2022 \times 2024$  convolution matrix:

$$R = \begin{pmatrix} -1 & 0 & 1 & & & \\ & -1 & 0 & 1 & & \\ & & -1 & 0 & 1 & \\ & & & \ddots & \ddots & \ddots \end{pmatrix}$$

Putting these together, we have

$$Y = CXX^T \iff \text{vec}(Y) = (R \otimes C)\text{vec}(X),$$

hence  $M = R \otimes C$ .

4. A reasonably smart way to compute this convolution is to calculate each element of  $Y$  by the formula  $Y_{ij} = -X_{i-1,j-1} + X_{i-1,j+1} - 2X_{i,j-1} + 2X_{i,j+1} - X_{i+1,j-1} + X_{i+1,j+1}$ , which costs 5 additions and 2 multiplications per output. If we let  $N = 2022^2$  be the number of outputs, this is  $\Theta(N)$  computational cost, with minimal storage ( $\Theta(N)$  for the output  $Y$ , and only a few numbers for the coefficients of the convolution).

In contrast, if we formed  $M$  explicitly as a dense matrix (i.e. storing all entries, whether they are zero or not) and computed  $\text{vec}(Y) = M\text{vec}(X)$ , then it would require  $\Theta(N^2)$  computational cost and storage, which is *millions of times* more expensive for an image of this size. In fact, just *storing*  $M$  as a dense array of 64-bit floating-point values would require  $2022^2 \times 2024^2 \times 8 \approx 1.34 \times 10^{14}$  bytes, or over *100 terabytes*, which is ridiculous.

Of course, there is a middle ground. One could store  $M$  in a “sparse matrix” data structure that only stores the nonzero entries. The cost then once again becomes  $\Theta(N)$ . But the efficiency will still be worse (by a constant factor), because of the complexity of looking up values in such a data structure doesn’t exploit the regularity of the convolution operation, versus the “matrix-free” approach of just applying the  $3 \times 3$  kernel to each pixel.

## Problem 6 (3+4+3+3 points)

Let  $f(A)$  be a function that maps  $m \times m$  matrices to  $m \times m$  matrices. Recall that its derivative  $f'(A)$  is a linear operator that maps any change  $\delta A$  in  $A$  to the corresponding change  $\delta f = f(A + \delta A) - f(A) \approx f'(A)[\delta A]$ , to first order in  $\delta A$ .

In this problem, you will study and prove a remarkable identity (Mathias, 1996): if  $f(A)$  is sufficiently smooth,<sup>1</sup> then for *any*  $\delta A$  (not necessarily small!) the following formula holds:

$$f\left(\underbrace{\begin{bmatrix} A & \delta A \\ & A \end{bmatrix}}_M\right) = \begin{bmatrix} f(A) & f'(A)[\delta A] \\ & f(A) \end{bmatrix}.$$

That is, one applies  $f$  to a  $2m \times 2m$  “block upper-triangular” matrix  $M$  (blank lower-left = zeros), and the desired derivative is in the upper-right  $m \times m$  corner of the result  $f(M)$ .

1. Check this identity numerically in Julia against a finite-difference approximation for  $f(A) = \exp(A)$  (the matrix exponential  $e^A$ , computed by `exp(A)` in Julia, or `expm` in Scipy or Matlab), for a random  $3 \times 3$   $A = \text{randn}(3,3)$  and a random small perturbation  $dA = \text{randn}(3,3) * 1e-8$ ; note that you can make the block matrix above by using `LinearAlgebra` followed by  $M = [A \ dA; \ 0I \ A]$ , and you can extract an upper-right corner by (e.g.) `M[1:3,4:6]`.
2. Prove the identity by explicit computation for the cases:  $f(A) = I$ ,  $f(A) = A$ ,  $f(A) = A^2$ , and  $f(A) = A^3$ . (Two of these are trivial! This is “bargain-basement induction”: do a few small examples and see the pattern.)
3. Prove the identity for  $f(A) = A^n$  for any  $n \geq 0$  by induction: assume it works for  $A^{n-1}$  and show using the product rule that it therefore must work for  $A^n$ . (You already proved the trivial  $n = 0$  base case in the previous part.)

*Remark:* Once it works for any  $A^n$ , it immediately follows that it works for any  $f(A)$  described by a Taylor series, such as  $\exp(A) = I + A + A^2/2 + A^3/6 + \dots + A^n/n! + \dots$ , since such a function is just a linear combination of  $A^n$  terms.

4. Prove the identity for  $f(A) = A^{-1}$  by explicit computation: since we know (from class) that  $f'(A)[\delta A] = -A^{-1}\delta A A^{-1}$ , plug this into the right-hand side of the formula above and show that it is the inverse of  $M$ : multiply by  $M$  and show you get  $I$ .

<sup>1</sup>The result is easiest to show when  $f(A)$  has a Taylor series (is “analytic”), and in fact you will do this below, but Higham (2008) shows that it remains true whenever  $f$  is  $2m - 1$  times differentiable, or even just differentiable if  $A$  is diagonalizable.

**Solution:**

1. See attached Julia notebook.
2. These three cases are:
  - (a)  $f(A) = A^0 = I$ : in this case  $f(M) = I$  ( $2m \times 2m$ ), and so the upper-right block is **zero**. This, of course, is the correct result  $d(I) = 0$ .
  - (b)  $f(A) = A$ : in this case,  $f(M) = M$ , and the upper-right block is  $\delta A$ . Again, this is the correct result:  $df = f'(A)[dA] = d(A) = dA$ , so  $f'(A)[\delta A] = \delta A$ .
  - (c)  $f(A) = A^2$ . In this case,

$$f(M) = M^2 = \begin{pmatrix} A & \delta A \\ & A \end{pmatrix} \begin{pmatrix} A & \delta A \\ & A \end{pmatrix} = \begin{pmatrix} A^2 & A\delta A + \delta A A \\ & A^2 \end{pmatrix}$$

by the usual “rows-times-columns” rule (which works for matrix *blocks* as well as for scalar elements). But then the upper-right block  $A\delta A + \delta A A$  is precisely  $f'(A)[\delta A]$  as derived in class by the product rule.

- (d)  $f(A) = A^3$ . Building off the previous part, we have

$$f(M) = M^3 = M^2 M = \begin{pmatrix} A^2 & A^2 \delta A + \delta A A \\ & A \end{pmatrix} \begin{pmatrix} A & \delta A \\ & A \end{pmatrix} = \begin{pmatrix} A^3 & A^2 \delta A + A \delta A A + \delta A A^2 \\ & A^3 \end{pmatrix}$$

again by the usual “rows-times-columns” rule. The upper-right block  $\delta A + A\delta A A + \delta A A^2$  is again  $f'(A)[\delta A]$  as derived in class, corresponding to  $d(A^3) = dA + A dA A + dA A^2$

3. For an inductive proof, we assume (for  $n > 0$ ) that the identity holds for  $n - 1$ , i.e. that:

$$M^{n-1} = \begin{pmatrix} A^{n-1} & (A^{n-1})'[\delta A] \\ & A^{n-1} \end{pmatrix}.$$

It then follows that

$$M^n = M^{n-1} M = \begin{pmatrix} A^{n-1} & (A^{n-1})'[\delta A] \\ & A^{n-1} \end{pmatrix} \begin{pmatrix} A & \delta A \\ & A \end{pmatrix} = \begin{pmatrix} A^n & A^{n-1} \delta A + (A^{n-1})'[\delta A] A \\ & A^n \end{pmatrix},$$

again by the usual “rows-times-columns” rule. But the upper-right block corresponds exactly to the product rule  $d(A^n) = d(A^{n-1} A) = A^{n-1} dA + d(A^{n-1}) A$ , so it is indeed  $(A^n)'[\delta A]$  as desired.

As noted in the problem, the inductive base case  $n = 0$  was already shown in the previous part, so our result must now hold for all  $n \geq 0$ .

4. For  $f(A) = A^{-1}$ , we know from class that  $f'(A)[\delta A] = -A^{-1} \delta A A^{-1}$ . We now want to show that

$$M^{-1} = \begin{pmatrix} A^{-1} & -A^{-1} \delta A A^{-1} \\ & A^{-1} \end{pmatrix},$$

which we can establish by explicit multiplication with  $M$  (on either the left or right):

$$\begin{pmatrix} A^{-1} & -A^{-1} \delta A A^{-1} \\ & A^{-1} \end{pmatrix} \begin{pmatrix} A & \delta A \\ & A \end{pmatrix} = \begin{pmatrix} A^{-1} A & A^{-1} \delta A - A^{-1} \delta A A^{-1} A \\ & A^{-1} A \end{pmatrix} = I$$

as desired:  $f(M) = M^{-1}$  indeed must have the correct derivative  $-A^{-1} \delta A A^{-1}$  in the upper-right block and  $A^{-1}$  on the diagonal blocks.